



Computational effort across the hybrid LLM-ILASP pipeline

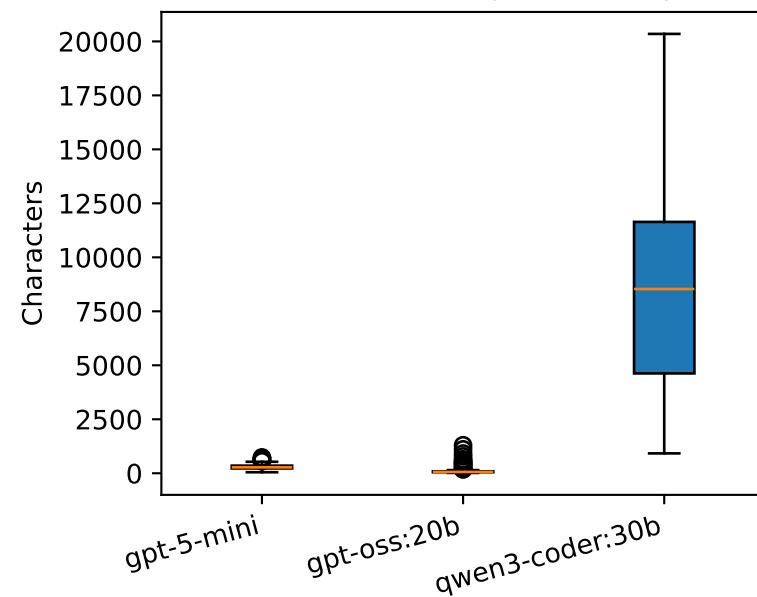
Thinking chars (Ollama)



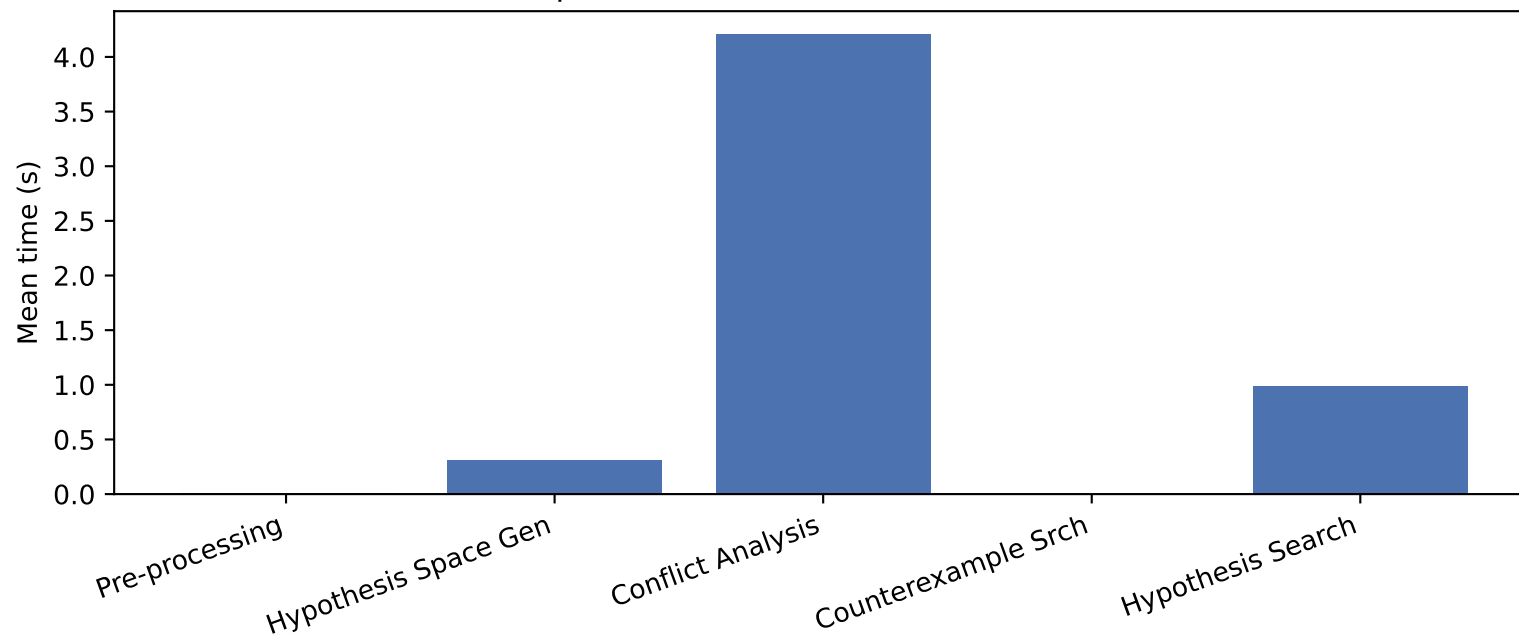
Thinking tokens (OpenAI)



Content chars (all models)



ILASP phase times (non-timeout benchmarks)



|H| vs CDILP iterations

