

LAMP: ASP Rule Induction and Optimization with LLMs

David Bunker
University of Potsdam

Abstract

Answer Set Programming (ASP) offers a compelling mechanism for knowledge representation and reasoning. However, developing such programs remains challenging. This work introduces the LLM to ASP Modeling Protocol (LAMP), a system by which an LLM iteratively generates and refines candidate ASP rules with each candidate verified by the CLINGO solver. LAMP is evaluated against the leading Inductive Logic Programming (ILP) system Inductive Learning of Answer Set Programs (ILASP), across 120 synthetic benchmarks of varying complexity. ILASP achieves 97.5% accuracy, with failures due to search space timeouts. Among LLMs, *gpt-5-mini* reaches 99.2% final accuracy, *gpt-oss:20b* 98.3%, and *qwen3-coder:30b* 64.2%. Although ILASP was on average faster, both *gpt-5-mini* and *gpt-oss:20b* solved all benchmarks on which ILASP timed out. Both LAMP and ILASP are also able to reduce literal counts and increase stratification to a similar degree. Finally, feature analysis is used to determine predictors of LLM success or failure.

1 Introduction

Knowledge representation and reasoning problems arise across many domains, including software and hardware customization, industrial manufacturing, service composition, and others. These problems are characterized by combinatorial constraints that determine which combinations of components are valid. Answer Set Programming (ASP) models such problems, offering clear semantics with expressive capabilities for constraints, defaults, and combinatorial reasoning.

However, writing ASP programs that are both correct and performant remains challenging, as modeling choices can significantly affect solver behavior. Recent progress in Large Language Models (LLMs) offers new opportunities to reduce this modeling burden. LLMs have shown strong capabilities in code generation and program synthesis, suggesting they may assist in deriving ASP encodings from high level natural language specifications [15].

The LLM to ASP Modeling Protocol (LAMP) is designed to guide LLMs in deducing and optimizing ASP rules from given instance facts and expected answer sets. LAMP provides a structured interaction protocol that enables refinement of ASP encodings, with the goal of improving both correctness and performance. Beyond proposing the system itself, LAMP can be used as a framework for evaluating the capabilities of different LLMs in ASP program generation.

Both LAMP and the leading Inductive Logic Programming (ILP) system Inductive Learning of Answer Set Programs (ILASP) [21] address the same underlying problem class. Given a set of instance facts and an expected answer set, find a program within a bounded hypothesis space that produces the correct answer set. ILASP solves this via conflict driven search under the Learning from Answer Sets (LAS) framework, while LAMP replaces symbolic search with LLM guided generation and CLINGO based iterative verification. As the expected input and output are the same, a controlled comparison can be made between symbolic ILP search and LLM based generation on the same class of tasks.

The relevant research questions of this work are the following.

RQ1. Is a structured approach for LLM assisted ASP modeling possible using the LAMP system, and under what conditions does it produce correct ASP encodings?

RQ2. How do LLMs and ILASP compare in terms of accuracy, speed, and program conciseness on ASP rule induction tasks?

RQ3. Which features of the ASP benchmark appear to predict LLM success or failure in LAMP?

While recent work has explored LLM and ASP integration through fine tuning [4], coupled planning pipelines [23], learning from answer sets [2], and solver feedback [27], LAMP is distinct in using LLMs to directly generate and iteratively refine ASP hypotheses verified by CLINGO, without requiring fine tuning or systematic hypothesis space search.

2 Background

This section presents the core background underpinning LAMP including ASP, ILASP, and LLMs.

2.1 Answer Set Programming

ASP is a declarative programming paradigm based on stable model semantics [11], well suited to combinatorial search and knowledge intensive reasoning tasks. An ASP program consists of rules over atoms and literals, and its solutions are the stable models, referred to as *answer sets*, of that program. The following introduces the syntactic primitives, rule structure, grounding procedure, and stable model semantics used throughout this paper.

Syntactic Primitives. A *term* is either a constant represented as a lower-case symbol or integer such as a , 0 or a variable represented as an uppercase symbol, such as X . A *predicate* p/k is a symbol of arity $k \geq 0$. Applying it to k terms yields an *atom* $p(t_1, \dots, t_k)$. A ground atom contains no variables. A *literal* is either an atom a , called a positive literal, or a *default negated* literal $\text{not } a$, where *not* denotes negation as failure and $\text{not } a$ holds whenever a cannot be derived.

Rules and Programs. An ASP program Π is a finite set of *rules* of the following form.

$$h \leftarrow b_1, \dots, b_m, \text{not } c_1, \dots, \text{not } c_n$$

Here h is the *head*, a classical atom or empty, each b_i is a positive body literal, and $\text{not } c_j$ denotes default negation of atom c_j . When $m = n = 0$ the rule is a *fact*. Every rule in an ASP program must be satisfied by a candidate answer set S . A rule is satisfied by S when either the body is not satisfied in S , or the head h is in S . A body is satisfied in S when every positive literal b_i is in S and every negated atom c_j is absent from S .

Variables and Grounding. Rules are written with first order variables, symbols such as X and Y that range over concrete terms rather than truth values, allowing compact expression over large domains. Prior to solving, a *grounder*, such as GRINGO, substitutes all relevant terms for each variable, producing an equivalent propositional ground program Π_g . A rule is *safe* if every variable occurring in it also appears in a positive body literal, which ensures that grounding is well defined.

Semantics. Let $\mathcal{A}$ denote the Herbrand base set of all ground atoms. Given a ground program Π and a set $S \subseteq \mathcal{A}$, the *Gelfond–Lifschitz reduct* Π^S is obtained by the following steps.

1. Remove every rule whose body contains $\text{not } c$ for some $c \in S$
2. Remove all negative literals from the bodies of remaining rules.

S is a *stable model*, referred to in ASP as an *answer set*, of Π if and only if S is the *minimal model* of the reduct Π^S . Because Π^S is negation free, it has a unique minimal model, the smallest set of atoms satisfying every rule. A program may have zero, one, or multiple answer sets, and solutions to the encoded problem correspond exactly to its answer sets. $AS(\Pi)$ is used to denote the set of answer sets of a program Π .

Modern ASP solvers, such as CLINGO [12], extend the core language with aggregates, choice rules, disjunctive heads, and optimization statements. LAMP uses CLINGO as its verification oracle but operates over a restricted subset of normal rules only, as defined in Section 3.1.

Solving Complexity. Solving non-ground normal ASP programs of bounded predicate arity is Σ_2^P -complete [6, 1], one level above **NP** in the polynomial hierarchy. Σ_2^P is the class of problems solvable by an **NP** machine with access to an **NP** oracle. The two levels of hardness arise because non-ground rules of bounded arity can compactly encode exponentially many ground instances, adding expressiveness beyond the search over candidate answer sets.

2.2 Learning from Answer Sets

ILASP [21] is an Inductive Logic Programming (ILP) solver that learns ASP programs from examples, defining the same task that LAMP approaches via LLM guided generation. A learning task is specified by a *hypothesis space* $\mathcal{H}$ constraining which rules may be learned, a *background knowledge* program B consisting of ASP rules that are given rather than learned, and a set of *learning examples* E specifying constraints on the answer sets that the learned hypothesis must produce. The following paragraphs define each component and describe the conflict driven search procedure ILASP uses to find an optimal hypothesis.

Hypothesis Space. The hypothesis space $\mathcal{H}$ is defined by a *mode bias*, which constrains the structure of rules that may appear in a learned hypothesis. A *mode declaration* is either a *head declaration* $\text{modeh}(a)$ or a *body declaration* $\text{modeb}(a)$, specifying which atoms may appear in rule heads and bodies respectively. Each atom in a mode declaration may contain *placeholder terms* of the form $\text{var}(t)$ (a variable of type t) or $\text{const}(t)$ (a constant of type t), which are instantiated during the search. A *hypothesis* $H \in \mathcal{H}$ is a set of ASP rules consistent with the mode bias.

Learning Examples. Examples are context dependent and constrain the answer sets of $B \cup H$. Each example specifies an inclusion set, an exclusion set, and a context program.

- **Positive examples** $e^+ \in E^+$: at least one answer set of $B \cup e_{\text{context}} \cup H$ must include all atoms in e_{included} and contain no atom in e_{excluded} .
- **Negative examples** $e^- \in E^-$: no answer set of $B \cup e_{\text{context}} \cup H$ may include all atoms in e_{included} while excluding all atoms in e_{excluded} .

To illustrate, suppose the goal is to learn which individuals are **allowed** entry, given facts about adults and minors as example context.

Listing 1: Simple ILASP learning task.

```
% Background knowledge B, could contain shared facts and rules

#constant(person, alice).
```

```

#constant(person, bob).
#constant(person, carol).

% Hypothesis space
#modeh(allowed(var(person))).
#modeb(adult(var(person))).
#modeb(minor(var(person))).

% Context dependent example: inclusion, exclusion, context
#pos(e1, {allowed(alice)}, {allowed(bob)}, {
    adult(alice). adult(carol). minor(bob).
}).

```

Background knowledge B may contain shared facts and rules, but here the ground facts are supplied as the example context instead. The context dependent positive example specifies an inclusion set, an exclusion set, and a context program. The inclusion set requires the hypothesis to derive `allowed(alice)`, ruling out an empty hypothesis. The exclusion set forbids `allowed(bob)`, ruling out the overly general rule `allowed(X)`. Together with the optimality criterion preferring shorter rules, they uniquely identify `allowed(X) :- adult(X)` as the intended hypothesis.

Learning Task. A *learning task* is a tuple $T = \langle B, \mathcal{H}, E^+, E^- \rangle$. A hypothesis $H \in \mathcal{H}$ is an *inductive solution* of T if it covers all examples, meaning for every $e^+ \in E^+$ some answer set of $B \cup e_{context}^+ \cup H$ satisfies e^+ , and for every $e^- \in E^-$ no answer set of $B \cup e_{context}^- \cup H$ satisfies e^- . When hypotheses are assigned a *cost*, typically the sum of rule lengths, the learning task seeks an *optimal* solution. This is the inductive solution of minimum cost.

Search Procedure. ILASP reduces the learning task to a sequence of ASP solving problems. The hypothesis space is encoded as an ASP metaprogram, and the search for a covering hypothesis is directed by a *conflict driven* procedure [21]. Candidate hypotheses are proposed and checked against all examples. Conflicts from examples not yet covered are then used to prune the space and guide subsequent iterations. This conflict driven approach scales substantially better than exhaustive enumeration of $\mathcal{H}$ [19].

Solving Complexity. Determining whether a satisfying hypothesis exists under the LAS framework is Σ_2^P -complete even for propositional programs [20], one level above the NP-completeness of propositional ASP solving. The additional level reflects the search over hypotheses $H \in \mathcal{H}$ layered on top of verifying that H covers all examples. This theoretical hardness motivates an alternative approach. LAMP replaces the hypothesis search with LLM guided generation, trading formal completeness for heuristic flexibility.




Figure 1: Overview of the LAMP system architecture.

2.3 Large Language Models for Code Generation

Large Language Models (LLMs) are neural language models trained on large amounts of text and source code, enabling strong performance on code generation tasks. Given a structured prompt, an LLM generates a candidate program token by token. Its output is nondeterministic and may require post processing or validation against a formal oracle.

LLMs have been shown to generate logic programs from natural language specifications [15, 22, 23] and to perform inductive logic programming when paired with a formal inference engine [9]. A key pattern across these systems is *iterative refinement* where the LLM generates a candidate, the candidate is verified symbolically, and any mismatch is encoded as feedback for the next prompt. LAMP instantiates this pattern for ASP, using CLINGO as the verification oracle and CLINGO derived feedback to guide successive generation attempts.

3 Methodology

The LAMP framework operates on a formally defined benchmark structure that constrains the hypothesis space and the target answer set, as illustrated in Figure 1. This section defines that benchmark framework, describes how benchmark instances are generated, and then details the LAMP interaction protocol and evaluation criteria.

3.1 Benchmark Framework

To keep the benchmarks tractable for both LLM generation and ILP search, the ASP used here is deliberately restricted. Programs consist only of unary *facts* of the form `predicate(term)` and *normal rules* of the form `new_predicate(X) :- predicate_0(X), [not] predicate_1(X), ...` which derive new predicates from initial instance facts, where `[not]` denotes optional default negation and rules may use multiple variables shared across body positions. This subset is expressive enough to require fairly complex rule induction while remaining within a reasonable hypothesis space.

The following is an example representing the model facts that terms "car" and "scooter" are both vehicles, represented with the predicate `vehicle`, and the scooter is slow moving, represented with the predicate `slow_moving`. It also has the rule that if a term is a vehicle and not slow moving it is highway possible, represented with the predicate `highway_possible`.

Listing 2: Vehicle example ASP program.

```
vehicle(scooter).
vehicle(car).
slow_moving(scooter).
highway_possible(X) :- vehicle(X), not slow_moving(X).
```

Running this results in the answer set $\{ \text{vehicle(scooter)}, \text{vehicle(car)}, \text{slow_moving(scooter)}, \text{highway_possible(car)} \}$. Given these facts and rules, `highway_possible(car)` can be deduced, indicating the car is highway possible and, due to default negation, the scooter is not highway possible.

This restricted ASP subset can now be formalized. A *Problem Profile* specifies how many predicates and terms exist, how many facts are given, and how large the rules may be. Together with the input facts and a target answer set, it defines a concrete benchmark that both LAMP and ILASP can solve.

Previous work on LLM based logic rule induction has characterized complexity by structural rule graph topology such as Chain, RDG (Rooted Directed Graph), DRDG (Disjunctive Rooted Directed Graph), and Mixed categories [9]. Since this work is concerned with correct answer set deduction and rule optimization rather than rule graph structure, complexity is instead quantified by seven numerical features. These include the number of possible predicates, terms, facts, rules, positive body literals per rule, negative body literals per rule, and the minimum number of derived atoms in the answer set. All generated programs are also limited to having exactly one answer set.

Definition 1 (Problem Profile). *A Problem Profile is a tuple*

$$\mathcal{C} = \langle \mathcal{D}, \mathcal{T}, N_F, R, L^+, L^-, A_{\min} \rangle$$

where

- $\mathcal{D}$ is a finite set of predicate symbols.
- $\mathcal{T}$ is a finite set of term constant symbols.
- $N_F \in \mathbb{N}$ is the number of instance facts.
- $R \in \mathbb{N}$ is the maximum number of rules.
- $L^+ \in \mathbb{N}$ is the maximum number of positive literals in each rule body.
- $L^- \in \mathbb{N}$ is the maximum number of negative literals in each rule body.
- $A_{\min} \in \mathbb{N}$ is the minimum derived atoms, the minimum number of atoms that must appear in the answer set beyond the given input facts.

From a Problem Profile a concrete input instance can be defined. Each instance pairs a set of ground facts and a target answer set with a bounded schema that defines the hypothesis space both LAMP and ILASP search over.

Definition 2 (Input Instance). *Given a Problem Profile stated as $\mathcal{C} = \langle \mathcal{D}, \mathcal{T}, N_F, R, L^+, L^-, A_{\min} \rangle$ as stated in Definition 1, an instance of $\mathcal{C}$ is a tuple*

$$\mathcal{M} = \langle F, A^*, \Phi \rangle$$

where

- F is the instance facts, a set of N_F ground atoms of the form $d(o)$ with predicate $d \in \mathcal{D}$ and term $o \in \mathcal{T}$ such as $d0(o0)$, so $|F| = N_F$.
- A^* is the target answer set, the desired unique stable model that any correct solution Π must produce. The atoms in $A^* \setminus F$ are the derived atoms, those produced by rules rather than given as facts. It must satisfy $F \subseteq A^*$ and $|A^* \setminus F| \geq A_{\min}$.
- Φ is the hypothesis space, the full set of concrete ASP rules from which Π is drawn. Each rule in Φ has the form

$$d_h(V_0) :- p_1(V_1), \dots, p_k(V_k), \text{ not } n_1(V'_1), \dots, \text{ not } n_m(V'_m).$$

where $d_h, p_i, n_j \in \mathcal{D}$, each V_i and V'_j is drawn from a finite set of logical variables and may be shared across positions, $k \leq L^+$ positive body literals, and $m \leq L^-$ negative body literals as bounded in Definition 1. All predicates are unary and all rules must be safe, meaning every variable in a rule appears in at least one positive body literal.

A candidate program Π is any subset $\Pi \subseteq \Phi$ with $|\Pi| \leq R$. Every rule in Π uses only predicates from $\mathcal{D}$ and respects the bounds L^+ and L^- . Π is a solution to $\mathcal{M}$ if and only if

$$AS(F \cup \Pi) = \{ A^* \}$$

where $AS(\cdot)$ denotes the set of answer sets of a program, meaning the unique answer set of $F \cup \Pi$ is exactly A^* .

The hypothesis space Φ corresponds directly to the mode bias of the associated ILASP learning task generated for each benchmark.

3.2 Dataset Generation

To instantiate benchmarks from the problem profiles defined above, synthetic data is generated by introducing new possible predicates and terms as needed. To keep the programs sufficiently complex, but manageable by the LLMs, only one stable model is expected and at least one atom must be deduced. The vehicle example above would be represented as shown below.

Listing 3: Benchmark ASP program.

```
d0(o0).
d0(o1).
d1(o1).
d2(X) :- d0(X), not d1(X).
```

This results in the answer set $\{ d0(o0), d0(o1), d1(o1), d2(o0) \}$, corresponding to the vehicle example answer set.

From each benchmark instance $\mathcal{M}$, a corresponding ILASP task file is generated. The predicates that appear in derived atoms $A^* \setminus F$ are declared as **#modeh** atoms, and all predicates in $\mathcal{D}$ are declared as **#modeb** atoms in both positive and negative forms. Since each benchmark has exactly one answer set, the derived predicates are uniquely determined by A^* , making the **#modeh** declarations both necessary and sufficient. No valid solution can require a head predicate not already present in $A^* \setminus F$. The target answer set A^* is encoded as a single context dependent **#pos** example with the instance facts F as context. The inclusion set contains the derived atoms $A^* \setminus F$, the exclusion set contains all ground atoms formed from derived predicates and terms in $\mathcal{T}$ that do not appear in A^* , and the context contains the instance facts. The corresponding ILASP task for the vehicle example above is shown below.

Listing 4: Corresponding ILASP task.

```
#constant(obj, o0).
#constant(obj, o1).

#modeh(d2(var(obj))).

#modeb(1, d0(var(obj))).
#modeb(1, d0(var(obj)), (negative)).
#modeb(1, d1(var(obj))).
#modeb(1, d1(var(obj)), (negative)).
#modeb(1, d2(var(obj))).
```

```
#modeb(1, d2(var(obj)), (negative)).

#pos(eg1, {
  d2(o0)
}, {
  d2(o1)
}, {
  d0(o0).
  d0(o1).
  d1(o1).
}).
```

This gives ILASP the same hypothesis space over which LAMP operates.

3.3 Prompting and Feedback Loop

The initial prompt instructs the LLM to reconstruct ASP rules that, combined with the given facts, yield the target stable model. It states the rule schema constraints from Φ including normal rules with bounded body literals and no new facts. It tells the LLM to prefer the minimal rule set and provides the instance facts F and target answer set A^* as concrete inputs. The prompt also contains a short worked example demonstrating the expected output format. If CLINGO reports a mismatch, a feedback message is appended to the conversation containing the facts, the expected stable model, and the computed answer sets from the candidate rules, asking the LLM to produce corrected rules. All LLMs were queried at the default temperature of 1 with each benchmark run once per model. This iterative interaction is formalized below.

Definition 3 (LAMP Interaction Protocol). *A LAMP Interaction Sequence for instance $\mathcal{M} = \langle F, A^*, \Phi \rangle$ is a sequence of pairs*

$$\langle (\Pi_0, \delta_0), (\Pi_1, \delta_1), \dots, (\Pi_n, \delta_n) \rangle$$

where each Π_i is a candidate program generated by the LLM given F , A^* , and all prior feedback, and δ_i is the feedback signal produced by running $F \cup \Pi_i$ through CLINGO:

$$\delta_i = \begin{cases} \textit{correct} & \text{if } AS(F \cup \Pi_i) = \{A^*\}, \\ \textit{unsatisfiable} & \text{if } AS(F \cup \Pi_i) = \emptyset, \\ \langle AS(F \cup \Pi_i), A^* \rangle & \text{otherwise.} \end{cases}$$

The sequence terminates when $\delta_n = \textit{correct}$ or a maximum iteration count is reached. The final output Π_n is accepted as the solution if it satisfies $\mathcal{M}$ as given in Definition 2.

3.4 Evaluation

LAMP and ILASP systems are evaluated by correctness of answer set, time to solve and optimality of ASP programs. The ASP programs generated are evaluated after being run via CLINGO. The correctness, number of rules, and how close it is to the expected stable model can then be assessed.

Beyond correctness, the systems are evaluated on whether the LLM has produced a more concise encoding than the original synthetic program, formalized as follows.

Definition 4 (Program Conciseness). *Given a solution Π to instance $\mathcal{M}$, its literal count is*

$$L(\Pi) = \sum_{r \in \Pi} (|pos(r)| + |neg(r)|),$$

where $pos(r)$ and $neg(r)$ are the positive and negative body literals of rule r , and $L(\Pi)$ is the total number of body literals across all rules. A solution Π is more concise than Π' if $L(\Pi) < L(\Pi')$.

Many other notions of optimality are applicable to ASP programs, including body literal and rule count, stratification, tightness, total grounded rules, and ASP solving time. This work reports both literal count $L(\Pi)$ and unique literal count, the number of distinct body literals across all rules in Π . While $L(\Pi)$ measures total deductive weight, unique literal count captures predicate vocabulary breadth, where a program that reuses few predicates across many rules is structurally simpler than one with many distinct dependencies, even at the same $L(\Pi)$. Stratification is reported separately.

4 Experiments

Experiments were conducted on the synthetic data with varying complexities as specified in Section 3.2. A variety of online and local open weight LLMs were selected for testing to get a broad view of capabilities.¹

4.1 Experimental Setup

A total of 120 synthetic ASP programs were created with varying features for testing. All experiments were conducted on an Apple M3 Max with 48 GB unified memory. The LLMs chosen included *gpt-5-mini*, which was run via OpenAI’s API platform, as well as open weight models *gpt-oss:20b* and *qwen3-coder:30b* which were run locally using Ollama. For *gpt-5-mini*, timing reflects server processing time reported by the OpenAI API, excluding network latency. All ASP solving and analysis was done using CLINGO.

¹Source code, benchmark data, and ILASP task files are available at <https://github.com/dbunker/lamp>.

ILASP version 4 was run on each benchmark with a timeout of 420 seconds, using the flags `-ml=3` and `-max-rule-length=5` which bound the maximum total literals in the learned program and the maximum length of any individual rule respectively. These runtime bounds restrict ILASP to a subset of the full hypothesis space Φ defined in Definition 2. Benchmarks not solved within this limit were recorded as failures. Each LAMP run was allowed a maximum of 3 iterations, one initial generation followed by up to 2 rounds of feedback verified by CLINGO, as diminishing returns were observed beyond this point. Both systems receive the same input instance $\mathcal{M}$ (Definition 2), where the target answer set A^* is provided directly to the LLM in the prompt and encoded as a `#pos` example for ILASP. The selected parameters used to construct problem profiles are listed in Table 1. The positive literal bound is derived as $L^+ = \text{Overall Literals} - L^-$. All combinations were used except where the negative literal count equals or exceeds the overall literal count, which would leave no positive literals.

Table 1: Parameters associated with Definition 1. Overall Literals is a generation constraint on the sum $L^+ + L^-$ and is not a profile parameter.

Parameter	Variations
Predicates	[6, 10]
Terms	[10]
Facts	[6, 9]
Rules	[11]
Overall Literals	[2, 3]
Negative Literals	[0, 1, 2]
Minimum Derived Atoms	[1, 3]

4.2 Results and Discussion

Overall results are shown in Table 2. Both *gpt-5-mini* and *gpt-oss:20b* perform very well, deducing the correct ASP rules to get the expected stable model in almost all cases. LAMP using *qwen3-coder:30b* performed significantly worse, reaching only 64.2% final accuracy (77/120), with 19.2% correct on the first attempt. Although *qwen3-coder:30b* is tuned for programming, it may not have been tuned for ASP, resulting in poorer performance.

ILASP achieved 97.5% accuracy (117/120), of which all three failures were timeouts, not reasoning errors, occurring on benchmarks with high negation counts and large numbers of derived atoms. LLMs in LAMP, *gpt-5-mini* and *gpt-oss:20b*, both solved all 3 of the benchmarks on which ILASP timed out, demonstrating that LLM guided search can succeed in some instances where ILP search cannot.

The iterative feedback loop contributes meaningfully for all models. First attempt accuracy was 90.8% for *gpt-5-mini* (109/120), 83.3% for *gpt-oss:20b* (100/120), and 19.2% for *qwen3-coder:30b* (23/120). Subsequent CLINGO verified feedback iterations recovered 10 additional cases for *gpt-5-mini*, 18 for *gpt-oss:20b*, and 54 for *qwen3-coder:30b*. In all cases where the correct stable model was produced, the generated rules never exactly matched those of the original synthetic program, as multiple rule sets can produce the same answer set. Given the sample size of 120 benchmarks and a single run per model at the default temperature of 1, differences of a few cases between systems should be interpreted with caution, as results may vary across runs and statistical significance may vary.

Table 2: Comparison of all systems including ILASP baseline. First attempt is the accuracy on the initial generation. Final accuracy is after up to 2 feedback iterations. Failures for all systems are either max iterations reached or exceeding the 420 s timeout.

System	First attempt		Final	Failures
<i>gpt-5-mini</i>	109/120 (90.8%)	119/120 (99.2%)		1
<i>gpt-oss:20b</i>	100/120 (83.3%)	118/120 (98.3%)		2
<i>qwen3-coder:30b</i>	23/120 (19.2%)	77/120 (64.2%)		43
ILASP	—	117/120 (97.5%)		3

Table 3 details the three benchmarks on which ILASP timed out (benchmark indices 7, 27, and 67). All three have elevated grounded rules and hypothesis space sizes relative to the dataset average (Table 4), indicating a larger search space as the cause of timeout. Benchmark 67 has a notably smaller hypothesis space than 7 and 27 yet still timed out, suggesting grounded rules and negation complexity are sufficient to cause timeout independent of hypothesis space size alone. Both *gpt-5-mini* and *gpt-oss:20b* solved all three, while *qwen3-coder:30b* failed all three.

Table 3: Outcomes for the 3 benchmarks where ILASP timed out (420 s limit). ✓ = correct stable model produced. × = failed.

Benchmark	<i>gpt-5-mini</i>	<i>gpt-oss:20b</i>	<i>qwen3-coder:30b</i>	ILASP
7	✓	✓	×	Timed out
27	✓	✓	×	Timed out
67	✓	✓	×	Timed out

Figure 2 shows the breakdown of benchmark outcomes per LLM across the three categories of correct on the first attempt (green), correct only after CLINGO verified feedback (gold), and failed across all iterations (red).

Table 4: Structural profile of the 3 ILASP timeout benchmarks vs. dataset average.

Benchmark	Grounded rules	Hypothesis space	Negative literals
7	27	19906	11
27	20	16960	11
67	27	2206	11
Dataset avg	13.28	3255.22	8.80




Figure 2: Stacked bar chart of benchmark outcomes per model showing first attempt success (green), feedback success (gold), and failed (red).

Figure 3 shows the per benchmark wall clock time distribution for each system as a box plot. *gpt-5-mini* is the fastest LLM with a median of 22.7 and a maximum of 199.2, reflecting the low latency of the API hosted model. *gpt-oss:20b* is slowest with median 43.2 and max 420.0 seconds, hitting the timeout on one benchmark, due to the larger locally run model size. ILASP has a highly bimodal distribution. The median is only 0.66 seconds as most benchmarks are trivially solved by conflict driven search, but three benchmarks hit the 420.0 second ceiling, pulling the maximum to that value. *qwen3-coder:30b* shows a relatively tight distribution with a median of 35.7 and a maximum of 112.8 seconds, consistent with the model completing its iteration budget in roughly similar time regardless of difficulty. Table 5 gives the aggregate statistics.




Figure 3: Per benchmark wall clock time distribution per system (seconds). Boxes show the interquartile range. Whiskers extend to the full range. ILASP failures are capped at the 420 s timeout.

Table 5: Wall clock time per system (seconds). LLM times accumulate all feedback iterations. ILASP timeouts are recorded as 420 s.

Model	Total (s)	Mean (s)	Median (s)	Max (s)
<i>gpt-5-mini</i>	3805.46	31.71	22.70	199.22
<i>gpt-oss:20b</i>	8630.91	71.92	43.22	420.00
<i>qwen3-coder:30b</i>	4577.38	38.14	35.72	112.77
ILASP	1917.36	15.98	0.66	420.00

Table 6 gives the mean output literal count $L(\Pi)$ (Definition 4) per system at original literal counts of 22 and 33, 11 rules $\times$ 2 or 3 body literals per rule, split by correctness. Correct solutions show a large reduction relative to the originals across all systems, averaging 3 to 6 literals against originals of 22 or 33. This is consistent with the synthetic benchmark structure, where generated programs may include rules not strictly necessary to produce the target answer set.

Table 6: Mean output literal count $L(\Pi)$ per system for benchmarks with 22 and 33 original body literals, split by correctness. - indicates no data.

Orig. literals		<i>gpt-5-mini</i>	<i>gpt-oss:20b</i>	<i>qwen3-coder:30b</i>	ILASP
22	Correct	4.81	5.15	5.85	3.67
	Incorrect	–	0.00	3.27	–
33	Correct	3.49	3.82	3.84	2.71
	Incorrect	14.00	–	3.43	–

Figures 4, 5, 6, and 7 compare the unique literal count of the original ASP programs against that of ILASP and the LLM generated rules for each LLM tested using LAMP. Dot size indicates the number of examples. In each figure, the left plot shows correct answer sets and the right shows first attempt failures, with recovered cases in gold and failed cases in red.

Figure 4 shows *gpt-5-mini* is generally able to greatly reduce the number of unique literals needed, though in a small number of cases it produces slightly more than the original. As shown in Figure 5, *gpt-oss:20b* had similar results, while *qwen3-coder:30b* (Figure 6) had mixed results. ILASP (Figure 7) is similarly able to reduce literal counts, though it fails to improve conciseness on the 3 benchmarks where it timed out.




Figure 4: Unique literal counts for *gpt-5-mini*.




Figure 5: Unique literal counts for *gpt-oss:20b*.




Figure 6: Unique literal counts for *qwen3-coder:30b*.




Figure 7: ILASP unique literal counts (timed out in red on right).

4.3 Stratification Analysis

An ASP program is *stratified* if no predicate depends negatively on itself directly or transitively. Stratification is a structural property that simplifies semantics and is often correlated with lower solver complexity. Since the LAMP benchmarks include programs with varying negation depths, it is informative to assess whether LLM generated programs preserve this property.

Table 7: Stratification of solved programs per system. Original is the 120 benchmark programs.

System	Stratified / Solved
Original	49 / 120
<i>gpt-5-mini</i>	119 / 119
<i>gpt-oss:20b</i>	117 / 118
<i>qwen3-coder:30b</i>	77 / 77
ILASP	117 / 117

Table 7 shows that both ILASP and the LLM based systems produce programs that are substantially more stratified than the original benchmark programs, of which only 49 of 120 were stratified. ILASP, being a formal inductive learner, produces stratified programs in 117 of its 117 correct solutions. Among the LLMs, *gpt-5-mini* stratifies all 119 of its correct programs, closely followed by *gpt-oss:20b* at 117 of 118. Even *qwen3-coder:30b*, which solved fewer benchmarks, produced stratified programs in all 77 of its correct solutions. The degree of stratification indicates both learned and generated programs tend toward syntactic simplicity that avoids the cyclic negation present in the synthetic benchmarks.

4.4 Feature Importance

To identify which benchmark features predict LLM failure, an XGBoost classifier [3] was trained on syntactic and structural features of the input programs using the *qwen3-coder:30b* results. The other two LLMs had too few failures for a meaningful classifier. Feature importance is measured by gain, the average improvement in loss from splits using that feature. Under five fold cross validation the classifier achieves 61.7% accuracy with AUC 0.619, below the 64.2% majority class baseline, so the importances shown in Figure 8 should be treated as weakly indicative only. The top features by gain, *Solution atoms*, *Possible terms*, and *Predicates*, all reflect problem size and vocabulary breadth, consistent with the intuition that larger answer sets and broader predicate sets increase the modeling burden on the LLM.




Figure 8: XGBoost feature importance of ASP features predicting LLM success.

5 Related Work

5.1 Inductive ASP Learning and LAS

Recent work has explored the integration of LLMs and ASP with several approaches studying how LLMs can learn from or reason over ASP directly. ILASP [21], the leading ILP system for Learning from Answer Sets, was combined with LLMs by Borroto et al. [2] which investigated question answering and inductive learning from answer sets, introducing the LLM2LAS system, which uses LLMs to construct ILASP learning tasks from natural language stories and then runs ILASP to induce the required commonsense rules.

A key distinction from LAMP is that LLM2LAS uses LLMs to *build the ILP task* and then defers to ILASP for search, whereas LAMP uses LLMs to *directly generate* the hypothesis, bypassing systematic hypothesis space search. Declarative knowledge distillation techniques have been proposed to transfer structured reasoning from LLMs into symbolic representations by Eiter et al. [7].

5.2 Expressivity Graded Evaluation

Gandarela et al. [9] demonstrated that LLMs can perform inductive logic programming, learning logic rules from background knowledge and positive and negative examples when combined with feedback from a formal inference engine such as Prolog. The authors introduced the RuDaS benchmark framework and a systematic methodology for evaluating LLM based theory induction graded by rule expressivity as Chain, RDG, DRDG, and Mixed categories. A key finding from that work is that LLMs are limited by their capacity to track long predicate dependency chains, directly relevant to future LAMP experiments and improvements.

Beyond ASP specific research, logic programming has been used as a tool to evaluate and scaffold LLM reasoning. Hybrid GOFAL LLM systems show how expert systems can be rebuilt using LLM generated logic rules [10], while Xu et al. [30] show programming based test evaluators systematically assess LLM reasoning reliability. These works support the use of formal solvers, as in LAMP, to ground and validate LLM outputs.

5.3 LLM Benchmarking on ASP

ASPBench [28] provides benchmarks specifically for grading LLM ASP induction. Ren et al. [26] benchmark 14 LLMs across ASP entailment, verification, and answer set computation tasks, finding that answer set computation remains a significant challenge for LLMs without scaffolding, suggesting that structured prompting and iterative solver feedback, as in LAMP, are likely necessary for reliable ASP induction. Schrader et al. [27] adopt a related solver-in-the-loop strategy, but use solver feedback to curate training data

for supervised fine tuning on natural language to ASP translation for logic puzzles, whereas LAMP uses solver feedback at inference time to guide rule induction from facts and a target answer set. LLM-ARC [16] adopts an actor critic framing in which an LLM generates logic and an automated reasoning critic evaluates it, and Kaur et al. [17] use ASP as a scaffold to improve LLM reasoning robustness under conformal prediction. Both pair LLMs with formal solvers in a pattern similar to LAMP, applied to general reasoning rather than rule induction.

5.4 Planning and NeuroSymbolic Pipelines

Planning provides another closely related line of research. Logical chain of thought tuning for symbolic planning by Verma et al. [29] and LLM symbolic planning pipelines without expert input by Huang et al. [14] highlight the benefits of keeping solvers such as ASP planners in the loop. Established ASP planning benchmarks, including Plasp developed by Dimopoulos et al. [5], offer realistic domains that could extend LAMP beyond product configuration. Neuro symbolic refinement frameworks such as PEIRCE further demonstrate how LLMs and formal reasoning systems can iteratively improve each other [24]. Hybrid LLM and ASP systems have also been applied to natural language understanding [25], commonsense reasoning [32], and collaborative conversational agents [31].

6 Future Work

Several directions for future work follow naturally. Iterative prompting could be improved by more explicit chain of thought reasoning, potentially evaluating many candidate rules across multiple instance fact sets per iteration, with CLINGO acting as a tool within an LLM based agent loop. Another approach would be from the specific perspective of solving complexity, focusing on solution space features such as vertex cover size, tree depth, feedback vertex set, clique width, tree width, among others [13]. Prompt optimization frameworks such as DSPy [18] could be used to automatically refine prompts and feedback strategies. Even without using LLMs fine tuned to ASP [4], the models are always improving, so simply updating the models may result in large improvements.

Repeated trials would strengthen statistical confidence in the reported accuracy differences.

The synthetic datasets could be extended to richer ASP constructs, including higher arity predicates, choice rules, disjunctive heads, and aggregates, as well as more complex dependency structures such as those proposed by Gandarela et al. [9] regarding logic program expressivity.

Evaluation could move beyond synthetic configuration tasks to existing ASP datasets, including ASPBench [28], OOASP models [8], real world ASP

encodings, planning benchmarks such as Plasp and PDDL to ASP translations [5], and domains like rail scheduling.

Deeper natural language interfaces to ASP could be explored, covering bidirectional translation between text and ground facts, rules, stable models, and multiple natural languages.

Finally, optimization remains a promising avenue, including other ways of transforming ASP into more efficient ASP, compiling ASP fragments to lower level code, such as C, and tighter integration with solver level optimizations, enabling LLMs to reason not only about correctness but also about grounding and solving performance.

7 Conclusion

This work proposed a structured LLM to ASP Modeling Protocol (LAMP) designed to support the deduction and optimization of ASP rules from instance facts and expected stable models. The results show this is effective for guiding LLMs toward correct and often optimized ASP encodings. By constraining the modeling task and providing iterative feedback via an ASP solver, this system reduces the modeling burden associated with ASP while preserving correctness.

The tested models exhibited clear performance differences, with *gpt-5-mini* achieving 99.2% final accuracy (90.8% on first attempt) and *gpt-oss:20b* achieving 98.3% (83.3% on first attempt), while *qwen3-coder:30b* reached only 64.2% (19.2% on first attempt) despite its programming focus. These results highlight that strong general code generation ability does not necessarily translate to proficiency in logic programming, underscoring the value of ASP specific evaluation frameworks.

The ILASP baseline achieved 97.5%, with the three failures being search space timeouts rather than reasoning errors. This shows LLMs do not follow the same classical ILP search. They navigate the hypothesis space differently, using token-level generation rather than conflict driven search, and can succeed where conflict driven search may not. The iterative feedback loop further contributed recovery of many cases for the three LLMs, confirming its value as a component of the LAMP pipeline. ILASP was faster on average but timed out on three benchmarks that LAMP solved, while both systems substantially reduced literal counts relative to the synthetic originals and produced more stratified programs than the benchmarks themselves. Feature analysis on *qwen3-coder:30b* suggested solution atom count, possible term count, and predicate count as the top predictors of LLM failure.

Beyond assisting users in ASP modeling, this protocol provides a system for probing LLM reasoning and logic programming capabilities. Future work includes extending to richer ASP constructs such as disjunction and choice rules, incorporating solver level performance metrics, and exploring

fine tuned or neuro symbolic hybrids to further improve robustness and scalability.

References

- [1] Manuel Bichler, Michael Morak, and Stefan Woltran. The power of non-ground rules in answer set programming, 2016.
- [2] Manuel Borroto, Katie Gallagher, Antonio Ielo, Irfan Kareem, Francesco Ricca, and Alessandra Russo. Question answering with llms and learning from answer sets, 2025.
- [3] Tianqi Chen and Carlos Guestrin. Xgboost: A scalable tree boosting system, 2016.
- [4] Erica Coppolillo, Francesco Calimeri, Giuseppe Manco, Simona Perri, and Francesco Ricca. Lasp: Fine-tuning large language models for answer set programming, 2024.
- [5] Yannis Dimopoulos, Martin Gebser, Patrick Lühne, Javier Romero, and Torsten Schaub. plasp 3: Towards effective asp planning, 2018.
- [6] Thomas Eiter, Michael Fink, Hans Tompits, and Stefan Woltran. Complexity results for answer set programming with bounded predicate arities and implications. *Annals of Mathematics and Artificial Intelligence*, 51(2-4):123–165, 2007.
- [7] Thomas Eiter, Jan Hadl, Nelson Higuera, and Johannes Oetsch. Declarative knowledge distillation from large language models for visual question answering datasets, 2024.
- [8] Andreas Falkner, Anna Ryabokon, Gottfried Schenner, and Kostyantyn Shchekotykhin. Ooasp: Connecting object-oriented and logic programming, 2015.
- [9] João Pedro Gandarela, Danilo S. Carvalho, and André Freitas. Inductive learning of logical theories with llms: An expressivity-graded analysis, 2025.
- [10] Eduardo C. Garrido-Merchán and Cristina Puente. Gofai meets generative ai: Development of expert systems by means of large language models, 2026.
- [11] Martin Gebser, Roland Kaminski, Benjamin Kaufmann, and Torsten Schaub. *Answer Set Solving in Practice*. Synthesis Lectures on Artificial Intelligence and Machine Learning. Morgan & Claypool Publishers, 2012.

- [12] Martin Gebser, Roland Kaminski, Benjamin Kaufmann, and Torsten Schaub. Clingo = asp + control: Preliminary report, 2014.
- [13] Markus Hecher and Rafael Kiesel. Extended version of: On the structural hardness of answer set programming: Can structure efficiently confine the power of disjunctions?, 2024.
- [14] Sukai Huang, Nir Lipovetzky, and Trevor Cohn. Planning in the dark: Llm-symbolic planning pipeline without experts, 2024.
- [15] Adam Ishay, Zhun Yang, and Joohyung Lee. Leveraging large language models to generate answer set programs, 2023.
- [16] Aditya Kalyanpur, Kailash Karthik Saravanakumar, Victor Barres, Jennifer Chu-Carroll, David Melville, and David Ferrucci. Llm-arc: Enhancing llms with an automated reasoning critic, 2024.
- [17] Navdeep Kaur, Lachlan McPheat, Alessandra Russo, Anthony G Cohn, and Pranava Madhyastha. An empirical study of conformal prediction in llm with asp scaffolds for robust reasoning, 2025.
- [18] Omar Khattab, Arnav Singhvi, Paridhi Maheshwari, Zhiyuan Zhang, Keshav Santhanam, Sri Vardhamanan, Saiful Haq, Ashutosh Sharma, Thomas T. Joshi, Hanna Moazam, Heather Miller, Matei Zaharia, and Christopher Potts. Dspy: Compiling declarative language model calls into self-improving pipelines, 2023.
- [19] Mark Law. Conflict-driven inductive logic programming, 2022.
- [20] Mark Law, Alessandra Russo, and Krysia Broda. The complexity and generality of learning answer set programs. *Artificial Intelligence*, 259:110–146, 2018.
- [21] Mark Law, Alessandra Russo, and Krysia Broda. The ilasp system for inductive learning of answer set programs, 2020.
- [22] Naiqi Li, Peiyuan Liu, Zheng Liu, Tao Dai, Yong Jiang, and Shu-Tao Xia. Logic-of-thought: Empowering large language models with logic programs for solving puzzles in natural language, 2025.
- [23] Xinrui Lin, Yangfan Wu, Huanyu Yang, Yu Zhang, Yanyong Zhang, and Jianmin Ji. Clmasp: Coupling large language models with answer set programming for robotic task planning, 2024.
- [24] Xin Quan, Marco Valentino, Danilo S. Carvalho, Dhairya Dalal, and André Freitas. Peirce: Unifying material and formal reasoning via llm-driven neuro-symbolic refinement, 2025.

- [25] Abhiramon Rajasekharan, Yankai Zeng, Parth Padalkar, and Gopal Gupta. Reliable natural language understanding with large language models and answer set programming, 2023.
- [26] Lin Ren, Guohui Xiao, Guilin Qi, Yishuai Geng, and Haohan Xue. Can llms solve asp problems? insights from a benchmarking study (extended version), 2025.
- [27] Timo Pierre Schrader, Lukas Lange, Tobias Kaminski, Simon Razniewski, and Annemarie Friedrich. A solver-in-the-loop framework for improving llms on answer set programming for logic puzzle solving, 2025.
- [28] Stefan Szeider. Asp-bench: From natural language to logic programs, 2026.
- [29] Pulkit Verma, Ngoc La, Anthony Favier, Swaroop Mishra, and Julie A. Shah. Teaching llms to plan: Logical chain-of-thought instruction tuning for symbolic planning, 2025.
- [30] Zihao Xu, Junchen Ding, Yiling Lou, Kun Zhang, Dong Gong, and Yuekang Li. Socrates or smartypants: Testing logic reasoning capabilities of large language models with logic programming-based test oracles, 2025.
- [31] Yankai Zeng and Gopal Gupta. Reliable collaborative conversational agent system based on llms and answer set programming, 2025.
- [32] Yankai Zeng, Abhiramon Rajashekharan, Kinjal Basu, Huaduo Wang, Joaquín Arias, and Gopal Gupta. A reliable common-sense reasoning socialbot built using llms and goal-directed asp, 2024.